

02443 Stochastic Simulation - Project 2

Project choice 2

Group 23:

Gunn Persdóttir Jacobsen - s175799

Julie Clausen - s180733

Søren Skjernaa - s223316

22. June 2023

Contents

0	Note	2
1	Introduction	3
2	Method	4
2.1	Simulation of the hospital queuing system	4
2.2	Optimizing bed distribution through Simulated Annealing	4
3	Hospital simulation without F^*-patients and ward	7
4	Hospital simulation with F^*-patients and ward	9
4.1	Initial analysis of the impact of creating ward F^*	9
4.2	Optimizing the bed distribution for a hospital with ward F^*	11
4.3	Comparing the found optimal solutions	14
4.4	Discussion	16
5	Sensitivity Analysis	17
5.1	Change in arrival distribution	17
5.2	Change in maximum capacity	20
6	Discussion/Conclusion	21
7	Appendix	23
7.1	R - code	23

0 Note

Group formation:

We tried to find a group on the discussion forum and by writing our contact information on all blackboards in the exercise rooms but did not get any response (likely due to the fact that we work in R, as most other groups seem to use Python). To ensure we finished the project on time we started solving it and wrote the report anyways.

Concerning the report:

All R-code used to produce this report's results is viewable in the appendix.

1 Introduction

In this project, our goal is to develop a simulation model that assesses the impact of implementing a specific distribution of bed resources across multiple wards within a hospital. The aim is to only address the problem from a technical standpoint, knowing that the decision-making process also involves organizational and political complexities.

We are considering a period where the hospital has been forced to create a new temporary ward (this ward is called F^*) due to a countrywide epidemic. Given that the hospital has limited resources, they are forced to reorganize using only staff and beds from the current inpatient wards. Currently the hospital has 5 different wards, $\mathcal{W} \in \{A, B, C, D, E\}$ each corresponding to 5 different patient types, represented by $\mathcal{P} \in \{A, B, C, D, E\}$. Every ward has a fixed bed capacity M_i . Patients of type $i \in \mathcal{P}$ arrive to ward $i \in \mathcal{W}$ with exponentially distributed inter-arrival time with rate λ_i . Furthermore, the patient stays at the hospital for an exponentially distributed amount of time with rate μ_i . The table in Figure 1 provides a summary of the parameters associated with each ward and patient.

Ward and patient type	Bed capacity	Arrivals per day (λ_i)	Mean length-of-stay ($1/\mu_i$)	Urgency points
A	55	14.5	2.9	7
B	40	11.0	4.0	5
C	30	8.0	4.5	2
D	20	6.5	1.4	10
E	20	5.0	3.9	5
F^*	<i>To be decided</i>	13.0	2.2	<i>Not relevant</i>

Figure 1: Parameters associated with each ward and patient type. Ward F^* denotes the new ward and the urgency points (column 5) reflect the "penalty" if a patient of type i is not admitted in Ward i .

In case a patient of type i is arriving at the hospital and the corresponding ward i is full, the patient is relocated to a different ward j . If the relocated ward also is full, then the patient is lost. Figure 2 presents a table of the relocation probabilities that patients of type i are being relocated to wards of type j (note that $i \neq j$).

$\mathcal{P} \backslash \mathcal{W}$	A	B	C	D	E	F^*
A	-	0.05	0.10	0.05	0.80	0.00
B	0.20	-	0.50	0.15	0.15	0.00
C	0.30	0.20	-	0.20	0.30	0.00
D	0.35	0.30	0.05	-	0.30	0.00
E	0.20	0.10	0.60	0.10	-	0.00
F^*	0.20	0.20	0.20	0.20	0.20	-

Figure 2: Probability, p_{ij} , of relocating a patient of type $i \in \mathcal{P}$ to an alternative Ward $j \in \mathcal{W}$. Includes the new Ward F^*

In the following we will also use $i = 1$ to denote patient type A and ward A, $i = 2$ to denote patient type B and ward B, and so on.

2 Method

We use a similar procedure for assessing the distribution of the beds effect on the number of relocated and lost patients in the hospital queuing system, for the scenario with and without F^* patients. In this section, we describe the main parts of the methods used.

As described in the introduction we use the convention that $i = 1$ corresponds to patient type A and ward type A, $i = 2$ corresponds to patient type B and ward type B, and so on.

2.1 Simulation of the hospital queuing system

We estimate the fraction of relocated and lost patients in the hospital queuing system, using the event-by-event principle. We simulate the arrival of $n = 10000$ patients, with an additional $k = 1000$ patients used for a burn-in period for the simulation.

For each patient, we use the following procedure (the procedure is illustrated without F^* -ward, but the necessary changes are easily assessed):

1. First, we sample an arrival time (time since the arrival of the previous patient), by sampling a single time from the $\exp(\sum_{i=1}^5 \lambda_i)$ distribution.
2. Next, we sample a patient type, by drawing from the 5-point distribution with probability $P(\text{Patient type} = i) = \frac{\lambda_i}{\sum_{i=1}^5 \lambda_i}$.
3. Using the sampled patient type i , we sample a time for the length of stay, by sampling from the $\exp(\mu_i)$ distribution.
4. Using the sampled patient type, we sample a potential relocation ward, by sampling from the 5-point distribution with probability $P(\text{Patient of type } i \text{ relocates to ward } j) = p_{ij}$.

We keep a record of each bed in the hospital, and the time left until the bed is free. When having sampled the four values for each patient, we attempt to assign the patient to a ward of the same type as the patient. If there is a free bed we adjust the time until the bed is free to match the patients length of stay. If the ward is full, we attempt to relocate the patient using the same bed procedure. If that ward is also full, we finally reject the patient from the system.

2.2 Optimizing bed distribution through Simulated Annealing

Using the urgency points given in table 1 we attempt to find an optimal distribution of the beds, which minimizes the penalty score for relocating patients.

Patient type	A	B	C	D	E	F^*
Urgency point	7	5	2	10	5	u_6

Table 1: Penalty for relocating (and or losing) patient of type i .

We implement a few different strategies for running simulated annealing and compare the obtained results. For the simulated annealing we again use a simulation with $n = 11000$ patients where the first 1000 are used for burn-in. Our simulated annealing process is described below:

0. We try two different methods for constructing the optimization problem:

- **Method 1:** We simulate 100 patient queues, which are reused for every step of the simulated annealing optimization process. When assessing the cost of a proposed bed distribution, we run the hospital simulation on each of the 100 patients and use the mean of the 100 costs as the cost for the proposed bed distribution.

The idea of reusing the same 100 patient queues is to avoid confounding potential improvements in the proposed bed distribution, with the randomness introduced by sampling new patient queues during optimization. We use the mean of 100 patient queues, to make the end result more robust towards different patient queues instead of only optimizing over one specific patient queue.

- **Method 2:** We run simulated annealing where we in each step, i.e. for each new proposed bed distribution, sample a new patient queue and find the cost for the proposed bed distribution on that patient queue.

The idea is that running simulated annealing with a new patient queue in each step should force the solution to become robust to the stochastic randomness in the distribution of the patient queue.

1. Initialize the bed distribution so that all wards have an equal amount of beds (with spare beds being given to ward A, then B, and so on).
2. Using the currently suggested bed distribution, we propose a random change to the bed distribution using the following procedure:
 - (a) Sample two wards $i, j = 1, 2, 3, 4, 5$ with equal probability.
 - (b) Sample a number of beds to swap $l = 1, 2, 3, 4, 5$ with probability $P(l = 1) = 0.6$, $P(l = 2) = 0.2$, $P(l = 3) = 0.1$, $P(l = 4) = 0.06$ and $P(l = 5) = 0.04$. We use decreasing probabilities, so that we in most cases only swap one bed, but in a few cases allow for the swapping of multiple beds, in order to sometimes get a large change in the bed distribution and possibly avoid being stuck in local minimas for the score.
 - (c) Remove l beds from ward i and add l beds to ward j .
 - (d) If all wards have at least 1 bed accept the proposed change, else reject it and return the unchanged bed distribution.
3. Calculate the cost of the proposed bed distribution, by rerunning the hospital simulation using the patient queue, and the proposed bed distribution. The cost is defined as $\text{Cost} = c + \sum_{i=1}^5 n_i u_i$, where u_i is the urgency points for patient type i , n_i is the number of

relocated and/or lost patients of type i , and c is the penalty for relocating and/or losing patients of type F^* . We have tried two possible methods for penalizing relocations of patients of type F^* .

- **Penalty type 1:** We use an urgency score of $u_6 = 10$ as the penalty for relocating (and or losing) patients of type F^* , so that $c = n_6 u_6$.
 - **Penalty type 2:** For each 0.1% above 5.0% of the patients of type F^* which are relocated, we penalize the score by 100. This means that if 5.7% of type F^* patients are relocated we penalize the score by 700. Thus $c = 100000(f_6 - 0.05)$, where f_6 is the fraction of relocated and/or lost patients of type F^* .
4. Given difference $\delta = \text{Cost}(\text{proposed bed distribution}) - \text{Cost}(\text{current bed distribution})$, we always accept the proposed bed distribution if $\delta \geq 0$. If $\delta < 0$ we accept with probability $e^{-\delta/T_i}$, where T_i is the "temperature" at step i of the simulated annealing process.

For the cooling of the temperature, we have used a temperature function defined by $T_i = \frac{T_0}{i}$ where the initial temperature $T_0 = 10000$ was found through experimentation. The result of our temperature function is seen in the following figures.

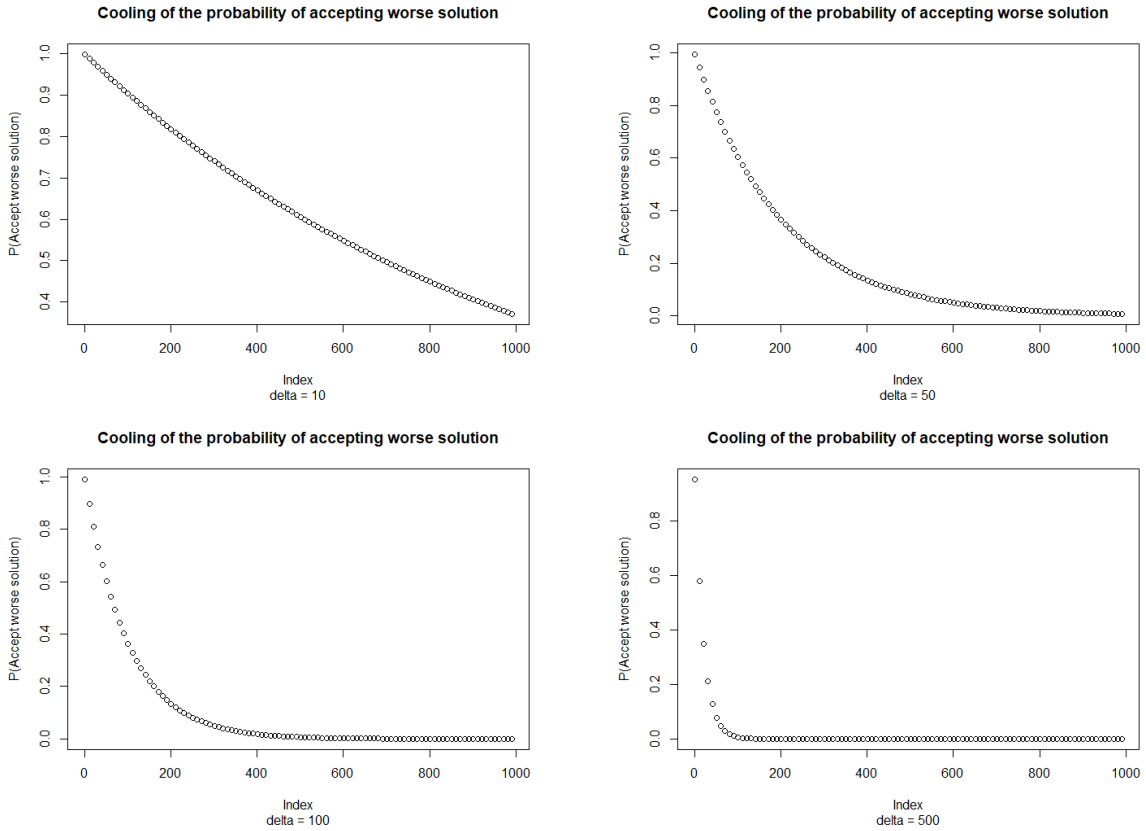


Figure 3: Probability of accepting worse solutions for $\delta = 10, 50, 100, 500$

We see that even towards the end of the simulated annealing process, we accept solutions that worsen the score by $\delta = 10$ with approximately 40% probability. For solutions, which

are worse by $\delta = 50$ the probability of accepting them drops to 40% after 200 steps, while it drops to 40% after 100 steps for $\delta = 100$. Thus, in the beginning, we are more open to large changes, while towards the end we do not accept solutions that change worsen the score by much.

3 Hospital simulation without F^* -patients and ward

We initiate the project by first, considering the starting case where we have 5 wards and 5 patient types and thereby exclude F^* at first. We then simulate the patient flow based on the given parameters in the hospital setting. Based on the simulation we conduct further calculations to study the probabilities for bed occupancy, admissions, relocations, and losses, and finally, we also consider the penalty scores. By doing so we are able to assess the impact of various factors on the efficiency of the hospital's ward management system and the allocation of beds in the different wards.

The parameters for this part are given in figure 1 and figure 2. When simulating $n = 10^4$ patients one time (with a burn-in period of 10^3) we get the following distribution of admitted, relocated, and lost patients.

Attribute	Ward A	Ward B	Ward C	Ward D	Ward E
N: patients	3254.000	2377.000	1795.000	1482.000	1092.000
N: admitted	3100.000	1909.000	1156.000	1423.000	795.000
N: relocated	97.000	321.000	532.000	48.000	215.000
N: lost	57.000	147.000	107.000	11.000	82.000

Table 2: Simulation Results

One notable finding in table 2 is that ward C has a relatively higher number of patient relocations compared to the other wards. This suggests that bed availability in ward C might have been insufficient, leading to the relocation of patients to different wards. But this also might make sende concidereing figure 2, since ward C has the highest mean length-of-stay per patient and only 30 beds. We also compute the penalty score, based on the urgency of patient types and the number of relocations and losses. The score for this particular simulation is equal to 6771. We use the score to assess the impact of different bed allocation scenarios or strategies within the simulated ward system later on.

Next, we run the simulation for $k = 10^2$ number of simulations and $n = 10^4$ number of simulated patients with a burn-in period of 10^3 . We run the simulation and calculate the probability of all beds being occupied upon arrival for each of the 5 wards.

Patient type	Mean	Variance	Lower 95% CI	Upper 95% CI
A	0.0459	0.0001	0.0437	0.0481
B	0.2112	0.0005	0.2070	0.2154
C	0.3473	0.0006	0.3426	0.3520
D	0.0289	0.0001	0.0267	0.0311
E	0.2887	0.0007	0.2833	0.2942

Table 3: Estimate and uncertainty for the probability of all beds being occupied upon arrival for each patient type.

In Table 3 we see the probabilities represented by the mean of each ward having no available beds for incoming patients. Higher probabilities suggest a higher likelihood of encountering a fully occupied ward upon arrival. Additionally, we have the variance which represents the spread or variability of the data around the mean. This means that a lower variance suggests that the values for a particular patient type are relatively close to the mean, while a higher variance indicates greater dispersion. In table 3, we see that the variance values are quite small, indicating that the data points for each patient type closely align with their respective means.




Figure 4: Histogram of the scores from the 100 simulations.

In Figure 4 the score distribution represents the distribution of penalty scores obtained from the simulations of patient allocation. The histogram divides the range of scores into 17 bins and displays the frequency or count of scores falling within each bin. When we calculate the mean, variance, and confidence interval we get the following values.

	Mean	Variance	Lower 95% CI	Upper 95% CI
Score	6883.33	433565.9	6752.68	7013.98

Table 4: Estimate and uncertainty for the score using the initial bed distribution.

The value of the mean represents the estimated score of the initial bed distribution. While the variance explains the variability of the scores around the mean. The confidence interval shows that this means that with a certain level of confidence of 95%, the true mean score is estimated to be within this range. Notice that the previous score we got with only one simulation is contained in this confidence interval.

4 Hospital simulation with F^* -patients and ward

In this section, we simulate the hospital queuing system, in the presence of patients of type F^* . The goal is to find an optimal distribution of the beds, where a maximum of 5% of patients of type F^* are being relocated (and/or lost).

In this section, we assume a total capacity of $k = 165$ beds across all wards of the hospital and we simulated the arrival of $n = 10000$ patients (with an additional 1000 patients used for burn-in of the simulation).

4.1 Initial analysis of the impact of creating ward F^*

As an initial test of the impact of including patients of type F^* in the simulation and creating ward F^* , we run 100 simulations of the hospital queuing system, where we use the bed distribution described in the introduction and move 20% of the beds from each of the wards A, B, C, D, E to ward F^* . Thus the bed distribution is given by:

Ward	A	B	C	D	E	F^*
Beds	44	32	24	16	16	33

Table 5: Bed distribution across the wards for the initial analysis.

Running the 100 simulations, we get the following probabilities for relocating (and/or losing) patients of each type:

Patient type	Mean	Variance	Lower 95% CI	Upper 95% CI
A	0.2392	0.0003	0.2355	0.2428
B	0.3885	0.0004	0.3843	0.3927
C	0.5415	0.0004	0.5375	0.5455
D	0.2169	0.0008	0.2114	0.2224
E	0.5473	0.0006	0.5426	0.5520
F^*	0.0608	0.0002	0.0582	0.0634

Table 6: Estimate and uncertainty for the probability of all beds being occupied upon arrival for each patient type.

We note that 6.08% of patients of type F^* are relocated (and/or lost) which does not meet the goal of only 5% of these patients being relocated. From the simulations we get the following estimate of the score when we use penalty type 1:

	Mean	Variance	Lower 95% CI	Upper 95% CI
Score	15532.44	591708.57	15379.81	15685.07

Table 7: Estimate and uncertainty for the score using the initial bed distribution.

The distribution of the scores (using penalty type 1) is furthermore seen in Figure 5




Figure 5: Histogram of the scores from the 100 simulations

4.2 Optimizing the bed distribution for a hospital with ward F^*

We use the different simulated annealing processes described in section 2.2 in order to find an optimal bed distribution w.r.t. the urgency points and the criteria that a maximum of 5% of patients of type F^* are not admitted to ward F^* .

Method 1, Penalty type 1: Figure 6, shows the development in the solution cost, when we ran simulated annealing for $k = 800$ steps reusing the same $l = 100$ patient queues in each step and used $u_6 = 10$ as the penalty for relocating patients of type F^* . We note, that improvements to the solution seem to have died out after 200 steps.

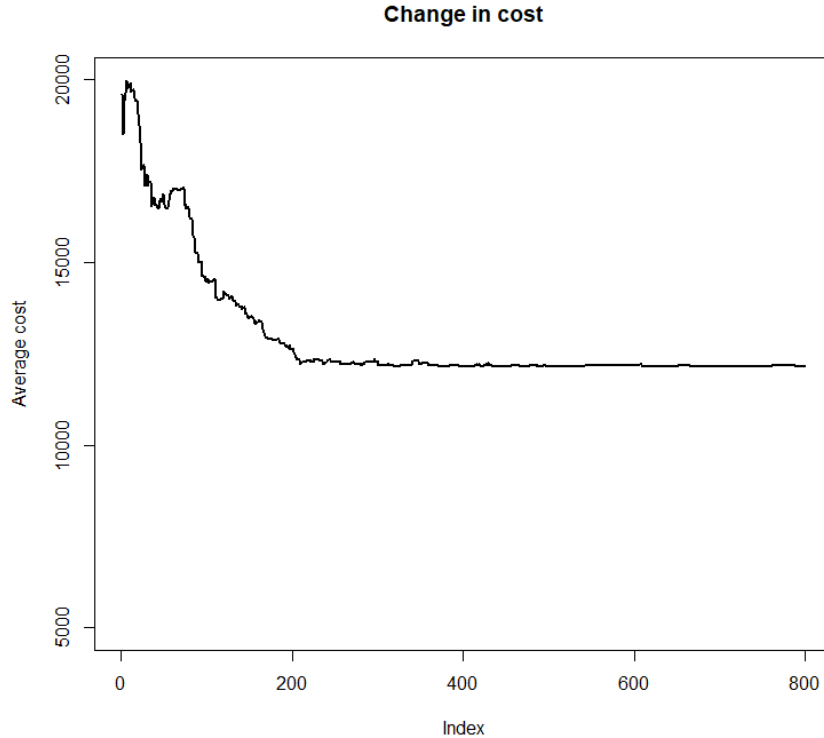


Figure 6: Development in the solution cost, when using method 1 and penalty type 1.

We found the minimal cost to be 12149.37 which was obtained when we used the following optimal bed distribution:

Ward	A	B	C	D	E	F^*
Beds	59	42	1	25	1	37

Table 8: Optimal bed distribution, when using method 1 and penalty type 1.

Method 1, Penalty type 2: Figure 7, shows the development in the solution cost, when we ran simulated annealing for $k = 400$ steps reusing the same $l = 100$ patient queues in each step and used a penalty of 100 for each 0.1% of patients of type F^* that was relocated above 5%. Again, we see that improvements to the solution largely seem to have died out after 200 steps and fully after 300 steps.

We found the minimal cost to be 11375.96 which was obtained when we used the following optimal bed distribution:

Ward	A	B	C	D	E	F^*
Beds	59	44	1	25	1	35

Table 9: Optimal bed distribution, when using method 1 and penalty type 2.

We see that there is no large difference between the solutions of penalty type 1 and 2.

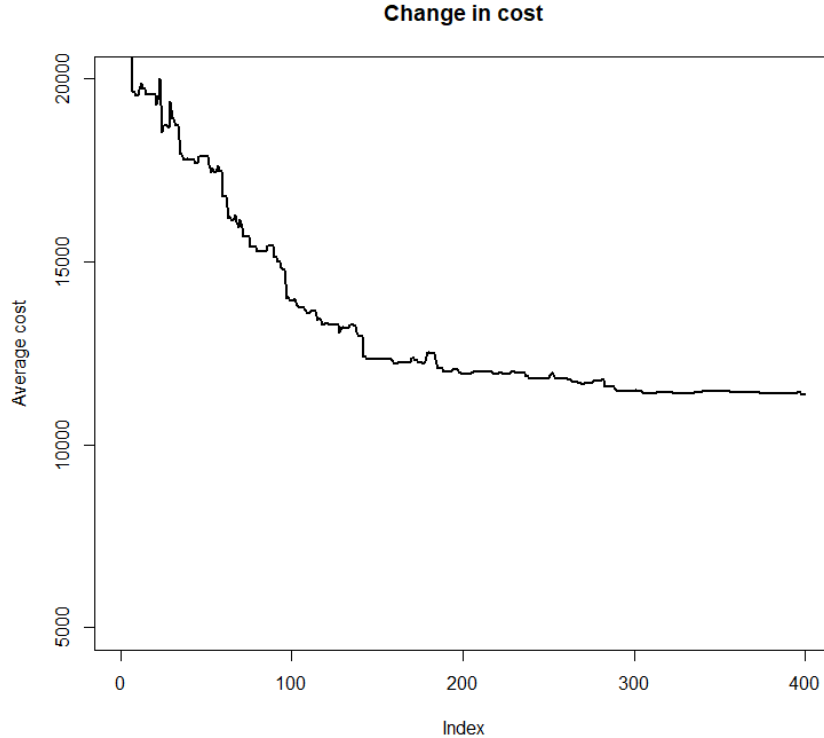


Figure 7: Development in the solution cost, when using method 1 and penalty type 2.

Method 2, Penalty type 1: Figure 8, shows the development in the solution cost, when we ran simulated annealing for $k = 400$ steps where we sampled a new patient queue in each step and used $u_6 = 10$ as the penalty for relocating patients of type F^* . We see that improvements to the solution largely seem to have died out after 150 steps.

We found the minimal cost to be 14201 which was obtained when we used the following optimal bed distribution:

Ward	A	B	C	D	E	F^*
Beds	37	25	22	27	19	35

Table 10: Optimal bed distribution, when using method 2 and penalty type 1.

Method 2, Penalty type 2: Figure 9, shows the development in the solution cost, when we ran simulated annealing for $k = 400$ steps where we sampled a new patient queue in each step and used a penalty of 100 for each 0.1% of patients of type F^* that was relocated above 5%. We see that improvements to the solution largely seem to have died out after 80 steps. When examining the proposed solution at each step we found that the small dip in the cost around step 390 is not due to a changed bed distribution, but a more optimal patient queue.

We found the minimal cost to be 13873 which was obtained when we used the following optimal bed distribution:

Ward	A	B	C	D	E	F^*
Beds	31	28	26	23	24	33

Table 11: Optimal bed distribution, when using method 2 and penalty type 2.

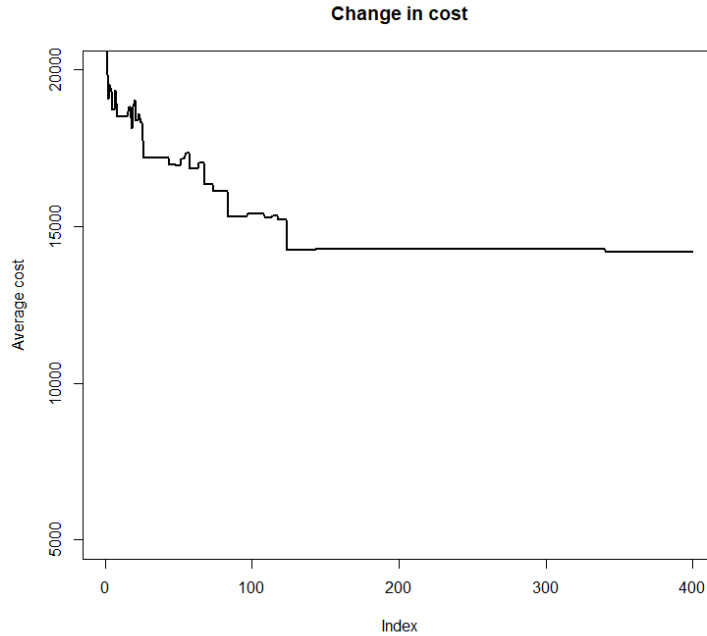


Figure 8: Development in the solution cost, when using method 2 and penalty type 1.




Figure 9: Development in the solution cost, when using method 2 and penalty type 2.

4.3 Comparing the found optimal solutions

While we do not see a large difference in the optimal bed distributions between penalty types 1 and 2, we do notice a large difference between methods 1 and 2. To compare the methods, we run 100 simulations for each of the four optimal solutions from tables 8, 9, 10 and 11 and calculate the obtained scores for each of the methods. The score is calculated using penalty type 1 for all four optimal solutions.

The distribution of the scores for each method is seen in figure 11, while table 12 shows the estimated scores with uncertainty

Score	Mean	Variance	Lower 95% CI	Upper 95% CI
Method 1, Penalty 1	12193.33	197319.23	12105.19	12281.47
Method 1, Penalty 2	12211.3	267870.7	12108.6	12314.0
Method 2, Penalty 1	16063.73	399628.32	15938.30	16189.16
Method 2, Penalty 2	17169.68	409433.01	17042.72	17296.64

Table 12: Estimate and uncertainty for the score using the initial bed distribution.

We see that the penalty type does not have a large influence on method 1. Furthermore, we see that method 1 outperforms method 2. When using method 2, we get worse score for both penalty type 1 and 2, than the scores we got from the initial test in table 7.

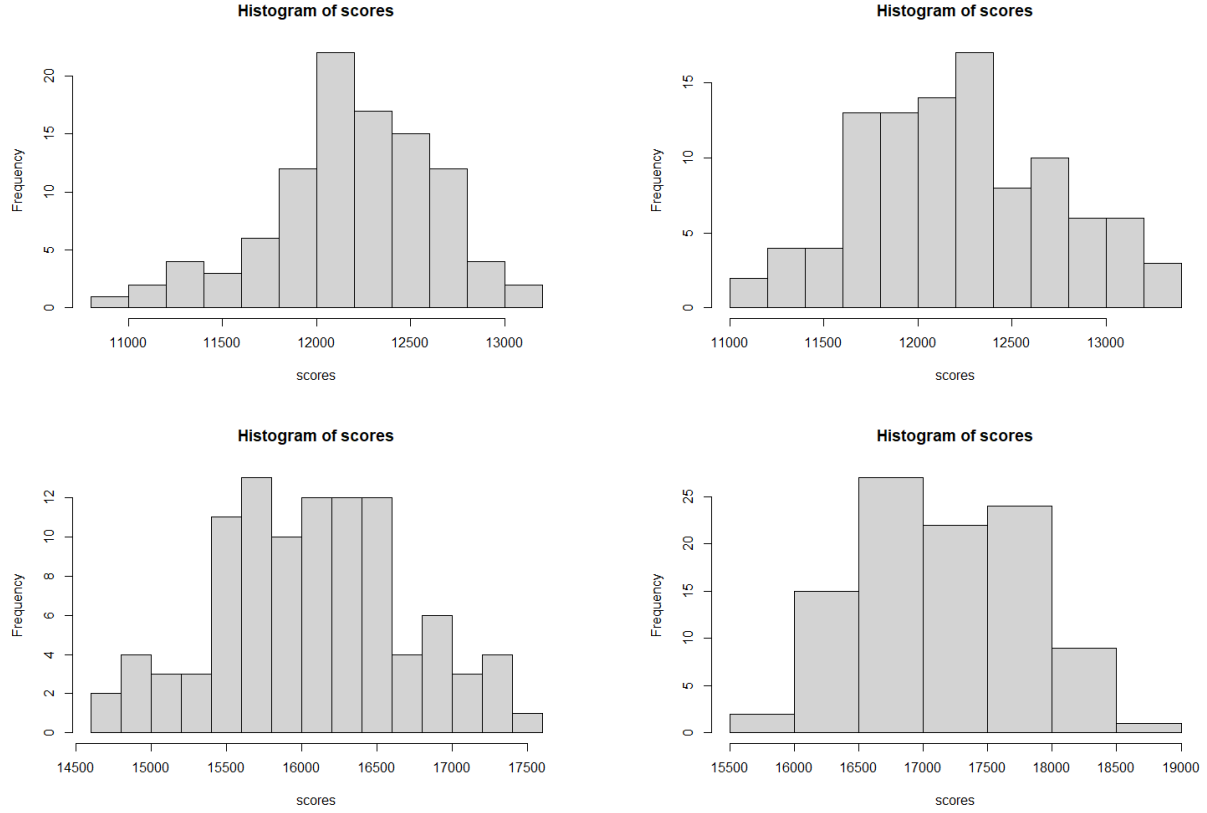


Figure 10: Histogram of the scores for each of the four methods. From left to right: Method 1 Penalty 1, Method 1 Penalty 2, Method 2 Penalty 1, Method 2 Penalty 2.

Table 13 below shows the probability of being relocated and/or lost for each patient type when using the optimal solution found from method 1 and penalty type 1. We note that patients of type F^* are relocated in less than 5% of the cases as we required from a solution.

Patient type	Mean	Variance	Lower 95% CI	Upper 95% CI
A	0.0976	0.0002	0.0946	0.1006
B	0.2304	0.0004	0.2264	0.2345
C	0.9815	0.0000	0.9809	0.9822
D	0.0666	0.0003	0.0632	0.0700
E	0.9708	0.0000	0.9696	0.9720
F^*	0.0296	0.0001	0.0280	0.0311

Table 13: Estimate and uncertainty for the probability of all beds being occupied upon arrival for each patient type for the optimal bed distribution found for method 1 and penalty type 1.

4.4 Discussion

We notice that for method 1 the optimal bed distribution was almost identical for penalty type 1 and 2, where as there was a small change for method 2. For both penalty type 1 and 2, method 2 failed to outperform our initial analysis where 20% of the beds from each ward were relocated to ward F^* . We suspect that this is more reflective of the way method 2 works, than the two penalty types.

When implementing method 2, we saw that the potential gains from proposing a new bed distribution was confounded with the increasing or decreasing cost of the new simulated patient queue in each step of the simulated annealing process. As such, improvements to the bed distribution quickly died out when using method 2. A reason for this may be that the algorithm is more likely to be stuck in a local minima which arose from a favorable patient queue in one step, and not because of a preferable bed distribution. This is avoided when using method 1, where the robustness towards changes in the patient queue comes from simulating and reusing multiple queues and finding a bed distribution that performs well on average over the queues.

We see from the optimal bed distributions that method 1 reduces ward C and E to only one bed each, and relocates the beds to wards A, B, D and F^* . This is in contrast to method 2, which also increases the number of beds in ward A and B but only slightly reduces ward C and E from the initial 27 beds in each ward. The result is that around 98% of all patients for ward C and E are relocated for method 1.

We see that method 1 outperforms method 2 in terms of the score. As ward C have the lowest urgency point and ward E is tied for second lowest together with ward B, method 1 optimizes the score by "ignoring" patients of type C and E in favour of reducing the number of relocations for the other patient types. As the mean length of stay is almost equal for patients B and E (4.0 and 3.9 respectively), and the arrival rate of type B is higher (11.0 compared to 5.0), the optimal bed distribution lowers the score by prioritizing beds for ward B before ward E.

We note that patients of type E have a 60% chance of being relocated to ward C. With both ward C and E having only one bed, patients of type E are likely to be completely lost, and not fill up beds in any of the other wards with higher priority patients. As there is no extra penalty for completely losing patients compared to only relocating them, method 1 obtains a better score by simply "turning away" patients of type E.

The bed distribution obtained from method 1, may reflect a real life scenario where in allocating resources for treating a new disease, we down prioritize certain patient groups and postpone their treatment. A more realistic model, which may avoid the strategy of "closing" certain wards, may be to add an extra penalty for completely losing patients, or to increase the length of stay for patients which are treated in a ward that does not match their patient type. We did not investigate such a simulation before having to hand in.

5 Sensitivity Analysis

5.1 Change in arrival distribution

In the sensitivity analysis of the system, we start out by testing how sensitive the system is to the length-of-stay distribution by replacing the exponential distribution with $\text{lognormal}(\mu, \sigma^2)$. For simplistic reasons, we have chosen to compare the system with the initial situation where the bed distribution was a given fixed distribution and where there was no ward or patients of type F^* .

To draw a length-of-stay from $\text{lognormal}(\mu, \sigma^2)$, we first find $\tilde{\mu}$ (the log mean) such the the mean length-of-stay is the same as in the initial set up. We do this by solving for $\tilde{\mu}_i$ in the following equation for $i = 1, \dots, 5$

$$\exp\left(\tilde{\mu}_i + \frac{\sigma^2}{2}\right) = \frac{1}{\mu_i}.$$

We'll be considering 3 different consecutively increasing variances, $\sigma_1^2 = 2/\tilde{\mu}_i^2$, $\sigma_2^2 = 3/\tilde{\mu}_i^2$ and $\sigma_3^2 = 4/\tilde{\mu}_i^2$. The following 3 tables show estimated mean and uncertainty for the probability of all beds being occupied upon arrival for each patient type, when the length-of-stay is drawn from $\text{lognormal}(\tilde{\mu}, \sigma_1^2)$, $\text{lognormal}(\tilde{\mu}, \sigma_2^2)$ and $\text{lognormal}(\tilde{\mu}, \sigma_3^2)$ respectively.

Patient type	Mean	Variance	Lower 95% CI	Upper 95% CI
A	0.0099	0.0001	0.0077	0.0121
B	0.1027	0.0017	0.0946	0.1108
C	0.1832	0.0026	0.1731	0.1932
D	0.0061	0.0001	0.0047	0.0075
E	0.1420	0.0035	0.1303	0.1536

Table 14: Estimate and uncertainty for the probability of all beds being occupied upon arrival for each patient type. The log-normal distribution is given $\sigma^2 = 2/\mu_i^2$

Patient type	Mean	Variance	Lower 95% CI	Upper 95% CI
A	0.4347	0.0019	0.4260	0.4435
B	0.5287	0.0023	0.5193	0.5382
C	0.6585	0.0023	0.6490	0.6681
D	0.3750	0.0032	0.3638	0.3862
E	0.6977	0.0017	0.6895	0.7059

Table 15: Estimate and uncertainty for the probability of all beds being occupied upon arrival for each patient type. The log-normal distribution is given $\sigma^2 = 3/\mu_i^2$

Patient type	Mean	Variance	Lower 95% CI	Upper 95% CI
A	0.6749	0.0010	0.6685	0.6813
B	0.7213	0.0010	0.7151	0.7276
C	0.8189	0.0008	0.8134	0.8244
D	0.6433	0.0026	0.6331	0.6534
E	0.8582	0.0006	0.8532	0.8632

Table 16: Estimate and uncertainty for the probability of all beds being occupied upon arrival for each patient type. The log-normal distribution is given $\sigma^2 = 4/\mu_i^2$

Starting by comparing table 14, 15 and 16 with table 3 we see that the estimated mean is somewhat lower in the first case when using σ_1^2 , but noticeable higher when using σ_2^2 and σ_3^2 . This is indicating that the system is somewhat sensitive of the variance of the time patients are staying at the hospital.

A substantial increase in the estimated mean for the probability of all beds being occupied results in a higher number of relocated and lost patients compared to the initial setup. Consequently, this leads to a larger penalty score, as depicted in the following tables.

	Mean	Variance	Lower 95% CI	Upper 95% CI
Score	3015.19	784608.3	2839.43	3190.95

Table 17: Estimate and uncertainty for the score using length-of-stay distribution as log-normal with variance $2/\mu_i^2$

	Mean	Variance	Lower 95% CI	Upper 95% CI
Score	27909.34	3312184	27548.22	28270.46

Table 18: Estimate and uncertainty for the score using length-of-stay distribution as log-normal with variance $3/\mu_i^2$

	Mean	Variance	Lower 95% CI	Upper 95% CI
Score	41003.78	1561715	40755.82	41251.74

Table 19: Estimate and uncertainty for the score using length-of-stay distribution as log-normal with variance $4/\mu_i^2$

The distribution of the scores from the 3 different log-normal distributions is furthermore seen in figure 11

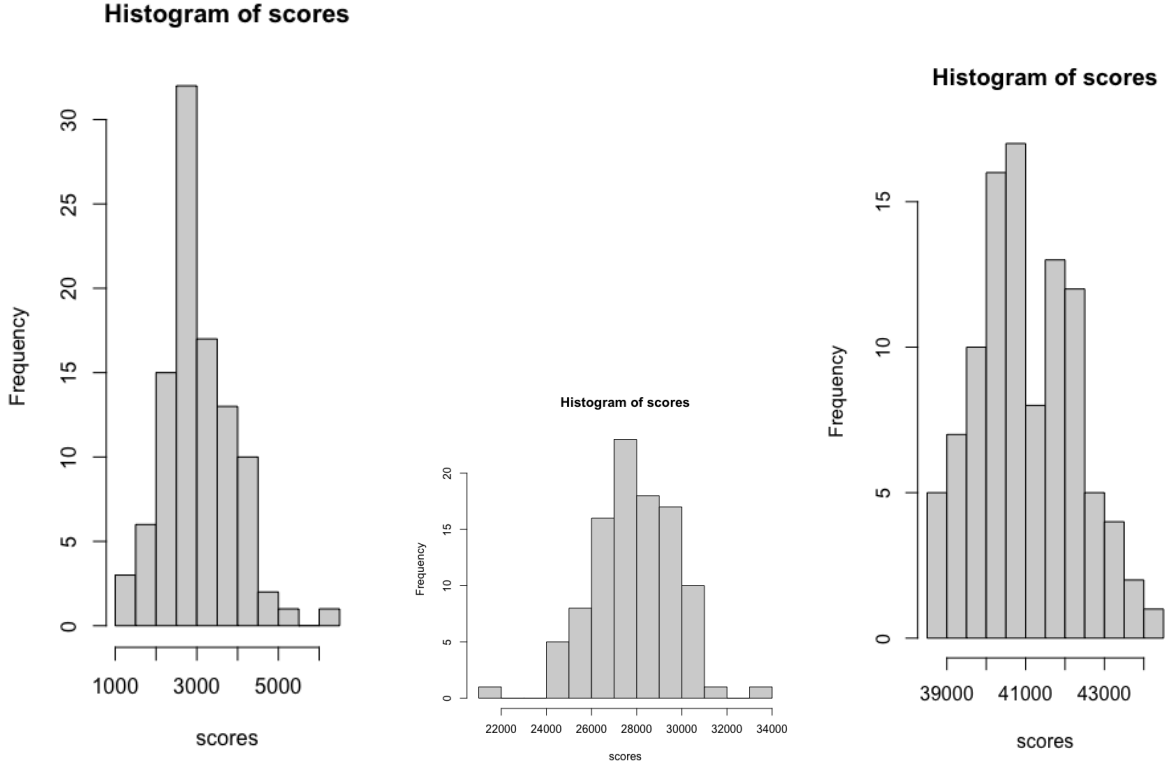


Figure 11: Histogram of the scores for each of the 3 log-normal distributions. From left to right: log-normal with variance $2/\mu_i^2$, log-normal with variance $3/\mu_i^2$, log-normal with variance $4/\mu_i^2$

5.2 Change in maximum capacity

Next we evaluate the hospital queuing systems sensitivity to a change in the total amount of beds. Again, we choose to focus on the original scenario without patients of type F^* . Furthermore, we distribute the beds among wards A, B, C, D and E following the same distribution as original, i.e we use the distribution given in table 20. This is done to show how stable the original distribution is to changes in the total bed capacity.

Ward	A	B	C	D	E
Fraction of total beds	0.34	0.25	0.18	0.12	0.12

Table 20: Bed distribution across the wards.

We run 100 simulations for each of the total bed capacities 150, 155, 160, \dots , 185. Averaging the probabilities of being relocated and/or lost for each patient type we get the results seen in figure 12.

Estimating the scores from the 100 simulations we get the estimates and 95% confidence intervals seen in table 21.

Total capacity	Mean	Variance	Lower 95% CI	Upper 95% CI
150	11216	600033	11063	11370
155	9702	462138	9567	9837
160	8130	400034	8005	8256
165	6755	295188	6647	6863
170	5625	283504	5519	5731
175	4802	250729	4702	4901
180	3961	174101	3878	4043
185	3433	177886	3350	3517

Table 21: Estimate and uncertainty for the scores for the different total capacities.

As we would expect the score and the probability of being relocated and/or lost drops as the total number of beds increase. The probability of being relocated drops to close to zero for patients A and D as the total bed capacity increases to 185. These two wards have the lowest mean length of stay (2.0 and 1.4 respectively). Thus, when wards A and D have 62 and 22 beds respectively almost no patients are relocated, and there is minimal gain from increasing the amount of beds on these two wards.

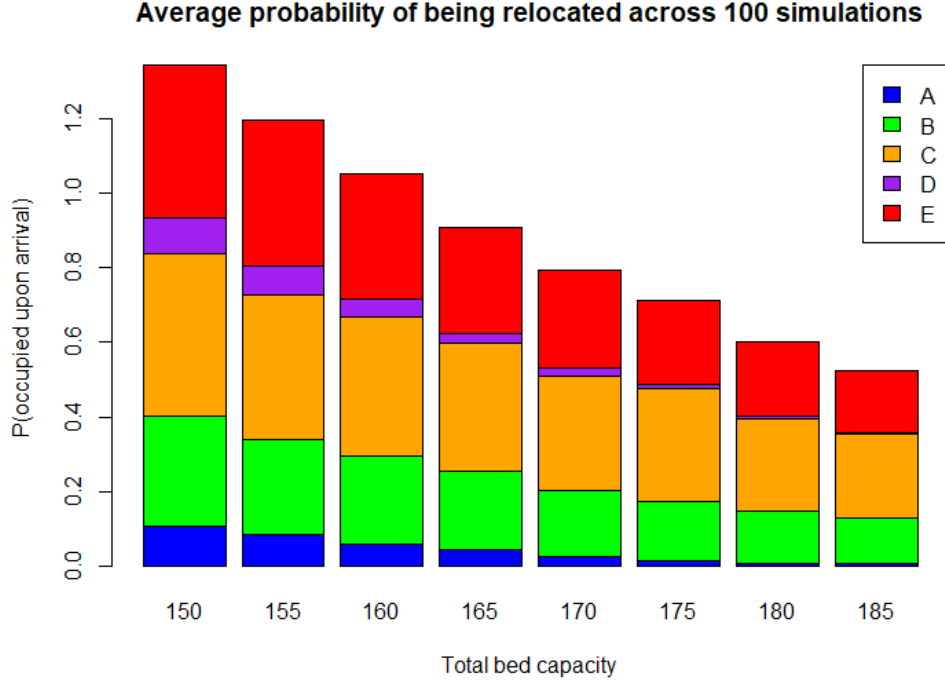


Figure 12: Probability of being relocated for each patient type.

6 Discussion/Conclusion

Based on the simulation of the hospital queuing system without the F^* ward, we obtained several important results. Firstly, the simulation provided insights into the distribution of admitted, relocated, and lost patients across the different wards and patient types. We observed variations in these numbers, indicating differences in the efficiency of bed allocation among the wards. For example, ward C had a higher number of patient relocations compared to the other wards, suggesting a potential issue with bed availability in that ward. On the other hand, ward B had a higher number of lost patients, indicating inefficiencies in accommodating patients in that particular ward.

Moving on to adding a new ward F^* and testing different strategies in this setting. We simulated the hospital queuing system with the inclusion of patients of type F^* and the goal of finding an optimal distribution of beds while ensuring a maximum of 5% of F^* patients are relocated or lost.

The initial analysis evaluates the impact of creating a new ward, F^* , by running 100 simulations with a bed distribution that reallocates 20% of beds from each existing ward to F^* . The simulations reveal that 6.08% of F^* patients are relocated or lost, exceeding the desired limit of 5%. To optimize the bed distribution, two methods are employed: Method 1, which reuses the same patient queues in each step, and Method 2, which samples a new patient queue in each step. Each method is tested with two penalty types: Penalty type 1,

where a fixed penalty is applied to relocating F^* patients, and Penalty type 2, where the penalty increases for each percentage point above the 5% relocation limit. Method 1 with Penalty type 1 yields a minimal cost of 12149.37 and an optimal bed distribution where beds are reallocated from wards C and E to other wards, reducing ward C and E to only one bed. Method 1 with Penalty type 2 results in a minimal cost of 11375.96, with a similar bed distribution. On the other hand, Method 2 with Penalty type 1 achieves a minimal cost of 14201, while Method 2 with Penalty type 2 attains a minimal cost of 13873.

Comparing the optimal solutions, it is observed that penalty type does not significantly impact Method 1, and Method 1 outperforms Method 2 regardless of the penalty type. However, all optimal solutions obtained from the methods have higher costs than the initial analysis without the inclusion of F^* ward. The probability of relocation or loss for each patient type in the optimal solution of Method 1 with Penalty type 1 is provided. Notably, F^* patients are relocated in less than 5% of cases, meeting the goal.

In the end Method 1 is more effective due to its ability to account for changes in patient queues, while Method 2 may get stuck in local minima. Furthermore, Method 1 redistributes beds from wards C and E to other wards, while Method 2 exhibits only minor changes. Leading to the analysis suggesting that Method 1 with Penalty type 1 provides the best results for optimizing the bed distribution while limiting the relocation of F^* patients.

Additionally, we examined the system's sensitivity to changes in the length-of-stay distribution in the sensitivity analysis. By replacing the exponential distribution with a lognormal distribution, we observed that the estimated mean probability of all beds being occupied upon arrival varied with different variances. The estimated mean generally increased as the variance of the lognormal distribution increased, indicating the sensitivity of the system to the variance of patients' hospital stays. This resulted in a higher number of relocated and lost patients, leading to larger penalty scores.

Furthermore, we evaluated the system's sensitivity to changes in the total bed capacity. We kept the bed distribution among wards consistent with the original scenario and conducted simulations for different total bed capacities. As the total number of beds increased, the probability of being relocated and/or lost decreased, and the penalty scores decreased accordingly. Wards with lower mean length of stay showed a smaller impact from increasing the bed capacity, as the probability of relocation decreased significantly for those wards.

7 Appendix

7.1 R - code

Simulation of hospital wards

```
1 # Preamble #####
2
3 ## File Description #####
4 #
5 #   Sren Skjernaa - s223316
6 #   22/06-2023
7 #
8 #   Stochastic Simulation
9 #   Project 2.2
10 #
11 #   Notes:
12 #       ...
13 #
14 #####
15
16 ## Clean up #####
17 rm(list = ls())
18 if(!is.null(dev.list())) dev.off()
19
20
21
22 # _ #####
23 # Ward simulation #####
24
25 ### Arguments:
26 #   - n          = Number of patients to simulate
27 #   - burn_in    = Burn_in period for the simulation
28 #   - capacity   = Capacities of the different wards
29 #   - lambda     = Arrival rate of patients of different types
30 #   - mu         = Length of stay rate of patients of different types
31 #   - P          = Probabilities of patient i transitioning to ward j
32 ###
33
34
35 ### Return:
36 #   - patients = Number of arrived patients of type i
37 #   - admitted = Number of admitted patients of type i to ward i
38 #   - relocated = Number of relocated patients of type i
39 #   - lost      = Number of lost patients of type i
40 ###
41
42
```

```

43 ward_sim <- function(n, burn_in, capacity, lambda, mu, P){
44
45     # number of different wards and types of patients
46     m <- length(capacity)
47
48     # initialize vectors for storage of results
49     patients <- numeric(m)
50     admitted <- numeric(m)
51     relocated <- numeric(m)
52     lost <- numeric(m)
53
54     # Initialize beds (time until bed is free)
55     max_capacity <- max(capacity)
56     temp <- c()
57     for (i in 1:m){
58         temp <- c(temp,
59                 rep(0, capacity[i]),
60                 rep(NA, max_capacity - capacity[i]))
61     }
62     beds <- matrix(temp, nrow = m, ncol = max_capacity, byrow = TRUE)
63
64
65     # Main loop of simulation - simulate journey of patient i
66     for (i in 1:(burn_in + n)){
67
68         # Sample time until one of the type of arrivals occur
69         arrival_time <- rexp(1, rate = sum(lambda))
70
71         # Reduce the time until bed is free
72         beds <- beds - arrival_time
73
74         # Sample type of patient that arrived
75         p_patient_type <- lambda / sum(lambda)
76         patient_type <- sample(1:m, 1, prob = p_patient_type)
77
78         # record patient
79         if (i > burn_in){
80
81             patients[patient_type] <- patients[patient_type] + 1
82         }
83
84         # Assign patient to ward of same type as patient
85         assigned_ward <- patient_type
86
87         # Sample length of stay for patient that arrived
88         length_of_stay <- rexp(1, rate = mu[patient_type])
89

```



```

90     # Check originally intended ward for free beds
91     bed_index <- which.min(beds[assigned_ward,])
92     bed_time <- beds[assigned_ward, bed_index]
93
94     if (bed_time <= 0){          # Assign patient to intended ward
95
96         beds[assigned_ward, bed_index] <- length_of_stay
97
98         # record admission
99         if (i > burn_in){
100
101             admitted[patient_type] <- admitted[patient_type] + 1
102         }
103
104     } else{                      # Try to relocate patient
105
106         # Sample a relocation
107         assigned_ward <- sample(1:m, 1, prob = P[patient_type,])
108
109         # Check relocated ward for free beds
110         bed_index <- which.min(beds[assigned_ward,])
111         bed_time <- beds[assigned_ward, bed_index]
112
113         if (bed_time <= 0){      # Assign patient to relocation ward
114
115             beds[assigned_ward, bed_index] <- length_of_stay
116
117             # record relocation
118             if (i > burn_in){
119
120                 relocated[patient_type] <- relocated[patient_type] + 1
121             }
122
123         } else{                  # Turn patient away
124
125             # record lost
126             if (i > burn_in){
127
128                 lost[patient_type] <- lost[patient_type] + 1
129             }
130         }
131     }
132 }
133
134
135
136 # Return results

```

```

137     result <- list("patients" = patients,
138                   "admitted" = admitted,
139                   "relocated" = relocated,
140                   "lost" = lost)
141
142     return(result)
143 }
144
145 # _ #####
146 # Simulation - without F #####
147
148 ## Test simulation - 1 run #####
149
150 ### Parameters -----
151 k <- 10^2                                # Number of simulations to run
152 n <- 10^4                                # Number of simulated patients
153 burn_in <- 10^3                          # Burn in period
154 wards <- c("A", "B", "C", "D", "E")     # Name of wards
155 capacity <- c(55, 40, 30, 20, 20)        # Ward bed capacity
156 lambda <- c(14.5, 11.0, 8.0, 6.5, 5.0)   # Ward arrivals per day
157 mu <- c(1/2.9, 1/4.0, 1/4.5, 1/1.4, 1/3.9) # Ward mean length of stay
158
159 P <- matrix(c(0, 0.05, 0.10, 0.05, 0.80,   # Probability patient i
160              0.20, 0, 0.50, 0.15, 0.15,    # relocates to ward j
161              0.30, 0.20, 0, 0.20, 0.30,
162              0.35, 0.30, 0.05, 0, 0.30,
163              0.20, 0.10, 0.60, 0.10, 0),
164             nrow = 5, byrow = TRUE)
165
166 urgency <- c(7, 5, 2, 10, 5)              # Penalty for relocating
167                                           # or losing patient of type i
168
169 ### Run simulation -----
170
171 # Run simulation
172 sim <- ward_sim(n, burn_in, capacity, lambda, mu, P)
173 patients <- sim$patients
174 admitted <- sim$admitted
175 relocated <- sim$relocated
176 lost <- sim$lost
177
178 # Find probability of all beds occupied upon arrival
179 occupied <- 1 - (admitted / patients)
180
181 # Calculate penalty score
182 score <- sum(urgency * (relocated + lost))
183

```

```

184
185 ### Results -----
186
187 # Format results in a table
188 temp <- matrix(c(occupied, patients, admitted, relocated, lost),
189               ncol = 5, byrow = TRUE,
190               dimnames = list(Attribute = c("P(occupied bed upon arrival)",
191                                           "N: patients",
192                                           "N: admitted",
193                                           "N: relocated",
194                                           "N: lost"),
195                               Ward = wards))
196 result_table <- as.table(temp)
197
198 # Print results
199 options(scipen = 999)
200 round(result_table, 3)
201 options(scipen = 0)
202
203 #Print score
204 score
205
206
207
208 ## Find distribution of score #####
209
210 ### Run simulations -----
211
212 # Initialize vector to store the scores
213 scores <- numeric(k)
214
215 # Run the simulations
216 for (i in 1:k){
217
218     sim <- ward_sim(n, burn_in, capacity, lambda, mu, P)
219
220     # Calculate score
221     scores[i] <- sum(urgency * (sim$relocated + sim$lost))
222 }
223
224 ### Show results -----
225
226 # Hisogram of scores
227 hist(scores, breaks = 15)
228
229 # Confidence intervals
230 mu <- mean(scores)

```

```

231 S2 <- var(scores)
232 t <- qt(0.975, df = k - 1)
233 mu
234 S2
235 c(mu - sqrt(S2 / k) * t, mu + sqrt(S2 / k) * t)
236
237
238
239 # _ #####
240 # Simulation - with F #####
241
242 ## Test simulation - 1 run #####
243
244 ### Parameters -----
245 k <- 10^2 # Number of simulations to run
246 n <- 10^4 # Number of simulated patients
247 burn_in <- 10^3 # Burn in period
248 wards <- c("A", "B", "C", "D", "E", "F") # Name of wards
249 # capacity <- c(28, 28, 28, 27, 27, 27) # Ward bed capacity
250 lambda <- c(14.5, 11.0, 8.0, 6.5, 5.0, 13.0) # Ward arrivals per day
251 mu <- c(1/2.9, 1/4.0, 1/4.5, 1/1.4, 1/3.9, 1/2.2) # Ward mean length of stay
252
253 P <- matrix(c(0, 0.05, 0.10, 0.05, 0.80, 0, # Probability patient i
254             0.20, 0, 0.50, 0.15, 0.15, 0, # relocates to ward j
255             0.30, 0.20, 0, 0.20, 0.30, 0,
256             0.35, 0.30, 0.05, 0, 0.30, 0,
257             0.20, 0.10, 0.60, 0.10, 0, 0,
258             0.20, 0.20, 0.20, 0.20, 0.20, 0),
259            nrow = 6, byrow = TRUE)
260
261 urgency <- c(7, 5, 2, 10, 5, 10) # Penalty for relocating
262 # or losing patient of type i
263
264 ### Run simulation -----
265
266 # Run simulation
267 sim <- ward_sim(n, burn_in, capacity, lambda, mu, P)
268 patients <- sim$patients
269 admitted <- sim$admitted
270 relocated <- sim$relocated
271 lost <- sim$lost
272
273 # Find probability of all beds occupied upon arrival
274 occupied <- 1 - (admitted / patients)
275
276 # Calculate penalty score
277 score <- sum(urgency * (relocated + lost))

```

```

278
279 ### Results -----
280
281 # Format results in a table
282 temp <- matrix(c(occupied, patients, admitted, relocated, lost),
283               ncol = 6, byrow = TRUE,
284               dimnames = list(Attribute = c("P(occupied bed upon arrival)",
285                                           "N: patients",
286                                           "N: admitted",
287                                           "N: relocated",
288                                           "N: lost"),
289                               Ward = wards))
290 result_table <- as.table(temp)
291
292 # Print results
293 options(scipen = 999)
294 round(result_table, 3)
295 options(scipen = 0)
296
297 #Print score
298 score
299
300 # _ #####
301 # Test of four optimal solutions #####
302
303 ## Parameters -----
304 k <- 10^2
305 n <- 10^4
306 burn_in <- 10^3
307 wards <- c("A", "B", "C", "D", "E", "F")
308 lambda <- c(14.5, 11.0, 8.0, 6.5, 5.0, 13.0)
309 mu <- c(1/2.9, 1/4.0, 1/4.5, 1/1.4, 1/3.9, 1/2.2)
310
311 P <- matrix(c(0, 0.05, 0.10, 0.05, 0.80, 0,
312              0.20, 0, 0.50, 0.15, 0.15, 0,
313              0.30, 0.20, 0, 0.20, 0.30, 0,
314              0.35, 0.30, 0.05, 0, 0.30, 0,
315              0.20, 0.10, 0.60, 0.10, 0, 0,
316              0.20, 0.20, 0.20, 0.20, 0.20, 0),
317            nrow = 6, byrow = TRUE)
318
319 urgency <- c(7, 5, 2, 10, 5, 10)
320
321 # capacity <- c(59, 44, 1, 24, 1, 36)
322 # capacity <- c(59, 44, 1, 25, 1, 35)
323 # capacity <- c(37, 25, 22, 27, 19, 35)
324 capacity <- c(31, 28, 26, 23, 24, 33)

```

```

325
326
327 ## Run simulations -----
328 p_occupied <- matrix(0, k, 6)
329 scores <- numeric(k)
330
331 for (i in 1:k){
332
333     # Simulation
334     sim <- ward_sim(n, burn_in, capacity, lambda, mu, P)
335
336     # P(relocated)
337     temp <- (sim$relocated + sim$lost) / sim$patients
338     p_occupied[i,] <- temp
339
340     # Score
341     scores[i] <- sum(urgency * (sim$relocated + sim$lost))
342 }
343
344 ## Show results -----
345 hist(scores, breaks = 10)
346
347 # Scores
348 m <- mean(scores)
349 S2 <- var(scores)
350 t <- qt(0.975, df = k-1)
351 CI <- c(m - sqrt(S2 / k) * t, m + sqrt(S2 / k) * t)
352 capacity
353 print(c(m, S2, CI))
354
355
356 # P(relocated)
357 for (i in 1:6){
358     mu <- mean(p_occupied[,i])
359     S2 <- var(p_occupied[,i])
360     t <- qt(0.975, df = k-1)
361     CI <- c(mu - sqrt(S2 / k) * t, mu + sqrt(S2 / k) * t)
362     print(round(c(i, mu, S2, CI),4))
363 }
364
365
366 # _ #####
367
368 # Sensitivity #####
369
370 ## Change in max capacity #####
371

```

```

372 ### Parameters -----
373 k <- 10^2
374 n <- 10^4
375 burn_in <- 10^3
376 wards <- c("A", "B", "C", "D", "E")
377 lambda <- c(14.5, 11.0, 8.0, 6.5, 5.0)
378 mu <- c(1/2.9, 1/4.0, 1/4.5, 1/1.4, 1/3.9)
379 P <- matrix(c(0, 0.05, 0.10, 0.05, 0.80,
380               0.20, 0, 0.50, 0.15, 0.15,
381               0.30, 0.20, 0, 0.20, 0.30,
382               0.35, 0.30, 0.05, 0, 0.30,
383               0.20, 0.10, 0.60, 0.10, 0),
384             nrow = 5, byrow = TRUE)
385 urgency <- c(7, 5, 2, 10, 5)
386
387 max_capacity <- seq(150, 185, 5)
388 original_capacity <- c(55, 40, 30, 20, 20)
389 bed_distribution <- original_capacity / sum(original_capacity)
390
391
392 ### Run simulations -----
393
394 m <- length(max_capacity)
395 p_occupied <- array(0, dim = c(k, 5, m))
396 scores <- matrix(0, k, m)
397
398 for (j in 1:m){
399
400   # Print iteration number
401   print(j)
402
403   # Divide beds among wards
404   capacity <- floor(bed_distribution * max_capacity[j])
405   temp <- c(rep(1, max_capacity[j] - sum(capacity)),
406             rep(0, 5 - (max_capacity[j] - sum(capacity))))
407   capacity <- capacity + temp
408
409   for (i in 1:k){
410
411     # Run simulation
412     sim <- ward_sim(n, burn_in, capacity, lambda, mu, P)
413
414     # P(relocated)
415     temp <- (sim$relocated + sim$lost) / sim$patients
416     p_occupied[i,,j] <- temp
417
418     # Score

```

```

419     scores[i,j] <- sum(urgency * (sim$relocated + sim$lost))
420   }
421 }
422
423
424 ### Results -----
425
426 # Find means
427 mean_scores <- colMeans(scores)
428 mean_occupied <- colMeans(p_occupied)
429
430
431 # bar chart
432 barplot(mean_occupied,
433         main = "Average probability of being relocated across 100 simulations",
434         xlab = "Total bed capacity", ylab = "P(occupied upon arrival)",
435         names.arg = max_capacity,
436         col = c("blue", "green", "orange", "purple", "red"))
437 legend("topright", legend = c("A", "B", "C", "D", "E"),
438       fill = c("blue", "green", "orange", "purple", "red"))
439
440 # Scores
441 plot(max_capacity, mean_scores, type = "p", pch = 16, col = "blue",
442      main = "Average score across 100 simulations",
443      xlab = "Total bed capacity", ylab = "Score")
444
445 # CI for scores
446 for (i in 1:m){
447   m <- mean(scores[,i])
448   S2 <- var(scores[,i])
449   t <- qt(0.975, df = k-1)
450   CI <- c(m - sqrt(S2 / k) * t, m + sqrt(S2 / k) * t)
451   print(round(c(max_capacity[i], m, S2, CI),0))
452 }
453
454
455 ### Show results -----
456 hist(scores, breaks = 10)
457
458 # Scores
459 m <- mean(scores)
460 S2 <- var(scores)
461 t <- qt(0.975, df = k-1)
462 CI <- c(m - sqrt(S2 / k) * t, m + sqrt(S2 / k) * t)
463 capacity
464 print(c(m, S2, CI))
465

```



```

466
467 # P(relocated)
468 for (i in 1:6){
469     mu <- mean(p_occupied[,i])
470     S2 <- var(p_occupied[,i])
471     t <- qt(0.975, df = k-1)
472     CI <- c(mu - sqrt(S2 / k) * t, mu + sqrt(S2 / k) * t)
473     print(round(c(i, mu, S2, CI),4))
474 }

```

Simulated Annealing

```

1  # Preamble #####
2  ## File Description #####
3  #
4  #   Sren Skjernaa - s223316
5  #   22/06-2023
6  #
7  #   Stochastic Simulation
8  #   Project 2.2
9  #
10 #   Notes:
11 #       ...
12 #
13 #####
14
15 ## Clean up #####
16 rm(list = ls())
17 if(!is.null(dev.list())) dev.off()
18
19
20
21 # _ #####
22 # Set up #####
23
24 ## Create test problems #####
25
26 ### Arguments:
27 #   - patients = Number of patients to simulate in each queue
28 #   - queues   = Number of patient queues to simulate
29 #   - lambda   = Arrival rate of patients of different types
30 #   - mu       = Length of stay rate of patients of different types
31 #   - P        = Probabilities of patient i transitioning to ward j
32 ###
33
34
35 ### Return:

```

```

36 #       - queue   = An array of dim(patients, 4, queue) storing the patient
37 #                   queues
38 ###
39
40 patient_queue <- function(patients, queues, lambda, mu, P){
41
42     # Initialize arrays to store patient queues
43     queue <- array(0, dim = c(patients, 4, queues))
44
45     # Create patient queues
46     for (i in 1:queues){
47         for (j in 1:patients){
48
49             # Number of different type of patients
50             m <- length(lambda)
51
52             # Draw arrival time
53             arrival_time <- rexp(1, rate = sum(lambda))
54
55             # Sample type of patient that arrived
56             p_patient_type <- lambda / sum(lambda)
57             patient_type <- sample(1:m, 1, prob = p_patient_type)
58
59             # Sample length of stay for patient that arrived
60             length_of_stay <- rexp(1, rate = mu[patient_type])
61
62             # Sample a relocation
63             relocation_ward <- sample(1:m, 1, prob = P[patient_type,])
64
65             # Store patient in queue
66             queue[j, ,i] <- c(arrival_time, patient_type,
67                               length_of_stay, relocation_ward)
68         }
69     }
70
71     # Return
72     return(queue)
73 }
74
75
76 ## Initial solution #####
77
78 ### Arguments:
79 #       - wards           = Number of different type of wards
80 #       - capacity        = Total number of beds across all wards
81 #       - initial_strategy = Either "uniform" or "random" initial distribution
82 ###

```

```

83
84
85 ### Return:
86 #       - bed_distribution = Initial distribution of beds across wards
87 ###
88
89 initial_distribution <- function(wards, capacity, initial_strategy){
90
91     if (initial_strategy == "random"){
92
93         # Generate random bed distribution across the wards
94         bed_distribution <- runif(wards)
95         bed_distribution <- bed_distribution / sum(bed_distribution)
96
97         # Scale the distribution so that it matches the max capacity
98         bed_distribution <- floor(bed_distribution * capacity)
99
100        # Add spare beds to ward with fewest beds
101        temp1 <- which.min(bed_distribution)
102        temp2 <- setdiff(1:wards, temp1)
103        bed_distribution[temp1] <- capacity - sum(bed_distribution[temp2])
104
105    } else if (initial_strategy == "uniform"){
106
107        # Generate uniform bed distribution accross the wards
108        bed_distribution <- rep(floor(capacity / wards), wards)
109        temp <- rep(1, capacity - sum(bed_distribution))
110        temp <- c(temp, rep(0, wards - length(temp)))
111        bed_distribution <- bed_distribution + temp
112    }
113
114    # Return
115    return(bed_distribution)
116 }
117
118
119 # _ #####
120 # Hospital simulation #####
121
122 ### Arguments:
123 #       - wards           = Number of wards and different types of patients
124 #       - queue           = Simulated patient queues
125 #       - bed_distribution = Current distribution of beds
126 #       - burn_in         = Burn in period of simulation
127 ###
128
129 ### Return:

```

```

130 # - patients = Number of arrived patients of type i
131 # - admitted = Number of admitted patients of type i to ward i
132 # - relocated = Number of relocated patients of type i
133 # - lost      = Number of lost patients of type i
134 ###
135
136 hospital_sim <- function(wards, queue, bed_distribution, burn_in){
137
138   # number of simulated queues and patients
139   m <- dim(queue)[3]
140   n <- dim(queue)[1]
141
142   # initialize vectors for storage of results
143   patients <- matrix(0, m, wards)
144   admitted <- matrix(0, m, wards)
145   relocated <- matrix(0, m, wards)
146   lost <- matrix(0, m, wards)
147
148   # Loop over all simulated queues
149   for (j in 1:m){
150
151     # Initialize time until bed is free counter
152     max_capacity <- max(bed_distribution)
153     temp <- c()
154     for (i in 1:wards){
155       temp <- c(temp,
156                 rep(0, bed_distribution[i]),
157                 rep(NA, max_capacity - bed_distribution[i]))
158     }
159     beds <- matrix(temp, nrow = wards, ncol = max_capacity, byrow = TRUE)
160
161     # Main loop of simulation - simulate journey of patient i
162     for (i in 1:n){
163
164       # Retrieve patients queue data
165       arrival_time <- queue[i, 1, j]
166       patient_type <- queue[i, 2, j]
167       length_of_stay <- queue[i, 3, j]
168       relocation_ward <- queue[i, 4, j]
169
170
171       # Reduce the time until bed is free
172       beds <- beds - arrival_time
173
174       # record patient
175       if (i > burn_in){
176

```

```

177     patients[j, patient_type] <- patients[j, patient_type] + 1
178 }
179
180 # Assign patient to ward of same type as patient
181 assigned_ward <- patient_type
182
183 # Check originally intended ward for free beds
184 bed_index <- which.min(beds[assigned_ward,])
185 bed_time <- beds[assigned_ward, bed_index]
186
187
188 if (bed_time <= 0){          # Assign patient to intended ward
189
190     beds[assigned_ward, bed_index] <- length_of_stay
191
192     # record admission
193     if (i > burn_in){
194
195         admitted[j, patient_type] <- admitted[j, patient_type] + 1
196     }
197
198 } else{                      # Try to relocate patient
199
200     # Assign patient to relocated ward
201     assigned_ward <- relocation_ward
202
203     # Check relocated ward for free beds
204     bed_index <- which.min(beds[assigned_ward,])
205     bed_time <- beds[assigned_ward, bed_index]
206
207
208     if (bed_time <= 0){      # Assign patient to relocation ward
209
210         beds[assigned_ward, bed_index] <- length_of_stay
211
212         # record relocation
213         if (i > burn_in){
214
215             relocated[j, patient_type] <-
216                 relocated[j, patient_type] + 1
217         }
218
219     } else{                  # Turn patient away
220
221         # record lost
222         if (i > burn_in){
223

```

```

224         lost[j, patient_type] <- lost[j, patient_type] + 1
225     }
226 }
227 }
228 }
229 }
230
231 # Return results
232 result <- list("patients" = patients,
233              "admitted" = admitted,
234              "relocated" = relocated,
235              "lost" = lost)
236 return(result)
237 }
238
239
240 # _ #####
241 # Evaluate solution #####
242
243 ## Cost of solution #####
244
245 ### Arguments:
246 #     - urgency = Penalty for relocating patients
247 #     - relocated = Number of patients that are relocated (and/or lost)
248 #     - patients = Number of patients of each type that arrives
249 #     - penalty = Penalty for relocating patients of type F (more than 5%)
250 ###
251
252
253 ### Return:
254 #     - score = Average score across all simulated queues
255 ###
256
257
258 cost <- function(urgency, relocated, patients, penalty){
259
260     # Number of simulated queues
261     m <- dim(relocated)[1]
262
263     # Initialize vector to store the scores
264     scores <- numeric(m)
265
266     # Loop over all simulated queues
267     for (i in 1:m){
268
269         # Calculate penalty for relocating patients of type F
270         if (penalty == 0){ # Penalty type 1

```

```

271     temp <- 0
272
273   } else{                                # Penalty type 2
274
275     temp <- (relocated[i, 6] / patients[i, 6]) * 1000
276     temp <- floor(max(0, temp - 50)) * penalty
277   }
278
279   # Calculate total score and save it
280   score <- sum(relocated[i,] * urgency) + temp
281   scores[i] <- score
282 }
283
284 # Return average score
285 result <- mean(scores)
286 return(result)
287 }
288
289
290
291 ## Plot solution #####
292
293 ### Arguments:
294 #   - current_distribution = Current bed distribution
295 #   - current_cost        = Current solution cost
296 #   - index               = Index of current solution
297 ###
298
299 plot_solution <- function(current_distribution, current_cost, index){
300
301   # Number of wards and total capacity
302   m <- length(current_distribution)
303   capacity <- sum(current_distribution)
304
305   # Labels for plotting
306   main_label <- paste("Capacity: ", capacity, " - Cost: ", current_cost)
307   sub_label <- paste("Index:", index)
308
309   # Plot solution
310   plot(1:m, current_distribution,
311        main = main_label, sub = sub_label, xlab = "Ward",
312        ylab = "Number of beds in wards",
313        type = "l", lwd = 2, col = "blue")
314 }
315
316
317 # _ #####

```

```

318 # Annealing strategy #####
319
320 ## Temperature cooling #####
321
322 temperature <- function(delta, index, T0){
323
324     # Calculate temperature (decreasing function of index)
325     # T = -log(index + 1)
326     # T <- 1 / sqrt(1 + index)
327     # T <- T0 * 0.90^index
328     T <- T0 / index
329
330     # Probability to accept a worse solution
331     p <- exp(- delta / T)
332     # p <- 0.25
333
334     return(p)
335 }
336
337 ### Plots of cooling #####
338 T0 <- 10^4
339 x <- seq(1, 10^3, 10)
340
341
342 plot(x, temperature(10, x, T0),
343      main = "Cooling of the probability of accepting worse solution",
344      sub = "delta = 10", xlab = "Index", ylab = "P(Accept worse solution)")
345
346 plot(x, temperature(50, x, T0),
347      main = "Cooling of the probability of accepting worse solution",
348      sub = "delta = 50", xlab = "Index", ylab = "P(Accept worse solution)")
349
350 plot(x, temperature(100, x, T0),
351      main = "Cooling of the probability of accepting worse solution",
352      sub = "delta = 100", xlab = "Index", ylab = "P(Accept worse solution)")
353
354 plot(x, temperature(500, x, T0),
355      main = "Cooling of the probability of accepting worse solution",
356      sub = "delta = 500", xlab = "Index", ylab = "P(Accept worse solution)")
357
358 plot(x, temperature(1000, x, T0),
359      main = "Cooling of the probability of accepting worse solution",
360      sub = "delta = 1000", xlab = "Index", ylab = "P(Accept worse solution)")
361
362
363
364 ## Change solution #####

```



```

365
366 ### Arguments:
367 #       - bed_distribution = Current distribution of beds
368 ###
369
370
371 ### Return:
372 #       - proposal_distribution = New proposed distribution of beds
373 ###
374
375 change_solution <- function(bed_distribution){
376
377     # Save old bed distribution
378     old_distribution <- bed_distribution
379
380     # Number of wards
381     m <- length(bed_distribution)
382
383     # Sample two random wards to take n beds from ward i and add to ward j
384     n <- sample(1:5, size = 1, prob = c(0.6, 0.2, 0.1, 0.06, 0.04))
385     swap <- sample(1:m, size = 2, replace = FALSE)
386     bed_distribution[swap[1]] <- bed_distribution[swap[1]] - n
387     bed_distribution[swap[2]] <- bed_distribution[swap[2]] + n
388
389     # Return new bed distribution
390     if (any(bed_distribution <= 0)){
391
392         return(old_distribution)
393
394     } else{
395
396         return(bed_distribution)
397     }
398 }
399
400
401 # _ #####
402 # Simulated Annealing #####
403
404 ## Main simulated annealing loop #####
405
406 ### Arguments:
407 #       - steps           = Steps to run the simulated annealing for
408 #       - patients        = Number of patients to simulate in each queue
409 #       - burn_in         = Burn in period to use
410 #       - queues          = Number of hospital queues to simulate
411 #       - capacity        = Maximum capacity across all wards

```

```

412 # - lambda      = Arrival rate of patients of different types
413 # - mu          = Length of stay rate of patients of different types
414 # - P           = Probabilities of patient i transitioning to ward j
415 # - urgency     = Penalty for relocating (or losing) patient i
416 # - initial_strategy = Uniform or random initial distribution of beds
417 # - T0          = Initial temperature
418 # - penalty     = Penalty for relocating patients of type F
419 # - plot_costs  = Logical variable controlling if development in
420 #                plots should be shown during optimization.
421 ###
422
423
424 ### Return:
425 # - solutions = Distribution of beds at each index
426 # - costs     = Costs of solution at each index
427 # - delta     = Improvement in solution for each proposal
428 # - patients  = Distribution of patient types in simulation
429 # - admitted = Number of admitted patients
430 # - relocated = Number of relocated patients
431 # - lost      = Number of lost patients
432 ###
433
434 sim_annealing <- function(steps,
435                           patients, burn_in, queues,
436                           capacity, lambda, mu, P, urgency,
437                           initial_strategy, T0, penalty,
438                           plot_costs){
439
440   # Parameters
441   total_patients <- patients + burn_in
442   wards <- length(lambda)
443
444   # Store costs and solution at each index
445   solutions <- matrix(0, steps, wards)
446   costs <- numeric(steps)
447   deltas <- numeric(steps)
448
449   # Create test problem (patient queues)
450   queue <- patient_queue(total_patients, queues, lambda, mu, P)
451
452   # Create initial solution (initial bed distribution)
453   bed_distribution <- initial_distribution(wards, capacity, initial_strategy)
454
455   # Find initial cost
456   sim <- hospital_sim(wards, queue, bed_distribution, burn_in)
457   total_relocations <- sim$relocated + sim$lost
458   current_cost <- cost(urgency, total_relocations, sim$patients, penalty)

```

```

459 # Simulated annealing loop
460 for (i in 1:steps){
461
462     # Loop counter (print to console)
463     print(i)
464
465     # Propose new solution by swapping adding bed from one ward to another
466     proposal_distribution <- change_solution(bed_distribution)
467
468     # Calculate cost of proposed bed_distribution
469     # queue <- patient_queue(total_patients, queues, lambda, mu, P)
470     sim <- hospital_sim(wards, queue, proposal_distribution, burn_in)
471     total_relocations <- sim$relocated + sim$lost
472     proposal_cost <- cost(urgency, total_relocations, sim$patients, penalty)
473
474
475     # Calculate change to cost
476     delta <- proposal_cost - current_cost
477
478
479     # Accept proposal if drop in cost: delta <= 0
480     if (delta <= 0){
481
482         bed_distribution <- proposal_distribution
483         current_cost <- proposal_cost
484
485         # Accept proposal with probability according to cooling scheme
486     } else if (runif(1) <= temperature(delta, i, T0)) {
487
488         bed_distribution <- proposal_distribution
489         current_cost <- proposal_cost
490     }
491
492     # Store current_cost
493     costs[i] <- current_cost
494     solutions[i,] <- bed_distribution
495     deltas[i] <- delta
496
497     # Plot development in cost during optimization
498     if (plot_costs){
499         plot(1:i, costs[1:i],
500             xlim = c(1, steps), ylim = c(5000, 20000),
501             type = "l", col = "black", lwd = 2,
502             main = "Change in cost", xlab = "Index", ylab = "Average cost")
503     }
504 }
505 }

```

```

506
507 # Run simulation a final time with the found optimum
508 sim <- hospital_sim(wards, queue, bed_distribution, burn_in)
509
510 # Return costs and solutions
511 result <- list("solutions" = solutions, "costs" = costs, "delta" = deltas,
512              "patients" = sim$patients, "admitted" = sim$admitted,
513              "relocated" = sim$relocated, "lost" = sim$lost)
514 return(result)
515 }
516
517
518 # _ #####
519 # Test Simulated Annealing #####
520
521 ## Parameters -----
522 steps <- 800
523 patients <- 10^4
524 burn_in <- 10^3
525 queues <- 100
526
527 capacity <- 165
528 lambda <- c(14.5, 11.0, 8.0, 6.5, 5.0, 13.0)
529 mu <- c(1/2.9, 1/4.0, 1/4.5, 1/1.4, 1/3.9, 1/2.2)
530 P <- matrix(c(0, 0.05, 0.10, 0.05, 0.80, 0,
531              0.20, 0, 0.50, 0.15, 0.15, 0,
532              0.30, 0.20, 0, 0.20, 0.30, 0,
533              0.35, 0.30, 0.05, 0, 0.30, 0,
534              0.20, 0.10, 0.60, 0.10, 0, 0,
535              0.20, 0.20, 0.20, 0.20, 0.20, 0),
536            nrow = 6, byrow = TRUE)
537 urgency <- c(7, 5, 2, 10, 5, 10)
538 penalty <- 0
539
540 initial_strategy <- "uniform"
541 T0 <- 10000
542 plot_costs <- TRUE
543
544 ## Run simulated annealing -----
545 sim <- sim_annealing(steps,
546                    patients, burn_in, queues,
547                    capacity, lambda, mu, P, urgency,
548                    initial_strategy, T0, penalty,
549                    plot_costs)
550
551 # Summary of results
552 sim

```

```
553 round((sim$relocated + sim$lost) / sim$patients, 3)
554 sim$solutions[steps,]
555 sim$costs[steps]
556 (sim$relocated + sim$lost) * urgency
557 sum((sim$relocated + sim$lost) * urgency)
558
559 test <- (sim$relocated + sim$lost) / sim$patients
560 mean(test[,1])
561 mean(test[,2])
562 mean(test[,3])
563 mean(test[,4])
564 mean(test[,5])
565 mean(test[,6])
```